

Master in High Performance Computing

Advanced Parallel Programming - MPI

Lab5 - MPI shared memory and Dynamic Process Management

Student: Tiago de Souza Oliveira

Summary of Optimization Strategies for `sumvalues.c` and Derivative Versions

1. Dynamic Process Management (DPM) and Singleton vs. Spawn Initialization

The optimization of `sumvalues.c` leveraged **MPI dynamic process management (DPM)** as defined in the *MPI 3.1 Standard*, primarily through the use of `MPI_Comm_spawn`. This allowed the application to start with a single parent process (singleton `MPI_Init`) and dynamically create an appropriate number of child processes based on the input workload dimensions (rows × columns).

While `MPI_Comm_spawn` introduces strong runtime flexibility and decouples parent-child computation responsibilities, it comes with practical caveats:

- It requires explicit support from the MPI implementation and underlying fabric (e.g., needing `I_MPI_SPAWN=on` for Intel MPI with OFI/MLX).
- Spawning is **heavier than launching with `mpirun -n N`**, and **less portable** under SLURM or tight HPC schedulers.

The comparison made with the **singleton `MPI_Init` with workload-based size check** and found that simpler designs often suffice when the universe size is fixed or known. This insight suggests using spawn-based DPM **only when runtime variability or recursive spawning is required**.

2. Shared Memory and MPI I/O Improvements

Two other dimensions of optimization were addressed:

- **Shared Memory (`sumvalues.c`):** Rewriting the application to use shared memory communicators (`MPI_Comm_split_type`) would enable intra-node memory access optimizations. However, this was deferred due to complexity vs. gain trade-offs in scenarios where file-based input dominates

bottlenecks.

- **MPI I/O (`sumvalues_mpio.c` / `sumvalues_mpio_dyn.c`):** It was adopted the collective file reading using `MPI_File_read_all` with subarray datatypes (`MPI_Type_create_subarray`). This provided deterministic and efficient partitioned I/O with less manual scattering of buffers. The `sumvalues_mpio_dyn.c` version additionally integrated dynamic process decisions using the same logic as the original `sumvalues.c`, but for collective I/O.

These I/O optimizations become **highly beneficial for large-scale matrix inputs** where traditional `fread` or scatter introduces overhead. However, for small matrices ($\leq 100 \times 100$), benefits were negligible or potentially outweighed by added complexity.

3. Key Feature Comparison & Future Recommendations

Version	Key Feature	Worth Further Investigation?
<code>sumvalues.c</code>	Static MPI setup, manual scatter, dynamic proc selection	✓ baseline, robust and simple
<code>sumvalues_spawn.c</code>	Runtime DPM via <code>MPI_Comm_spawn</code>	⚠ useful in dynamic or recursive apps, but fragile in HPC setups
<code>sumvalues_mpio.c</code>	Collective I/O, fixed workload size	✓ great for large files, scalable
<code>sumvalues_mpio_dyn.c</code>	Combined MPI I/O + dynamic workload sizing	✓ hybrid strategy, elegant balance

4. Conclusion

The evolution of `sumvalues.c` into multiple optimized forms demonstrates practical applications of modern MPI capabilities from the *MPI 3.1* standard. For most workloads in HPC environments, **singleton `MPI_Init` with controlled `mpirun` size and MPI I/O offers the best trade-off** in simplicity, portability, and performance. **Dynamic spawning should be reserved for adaptive workflows** or heterogeneous pipeline construction, where runtime flexibility is indispensable. Continuing exploration of MPI shared memory and `MPI_Win` may also complement this path when aiming for NUMA-aware, also OpenMP-SIMD acceleration.