

Assignment 4

Multiple matrix scaling with multiple outputs and overlapping

Problem definition

- Implement a variant of the CUDA version of the multiple matrix scaling presented for Assignment1 where the output are "n_scal" matrices, being "n_scal" the number of scaling factors.
- Each output matrix is the input multiplied by one of the factors
- You will probably need to modify the kernels: Create one kernel that scales the matrix for a certain factor and you must call the kernel as many times as scaling factors
- Overlap the transfer of output matrix for factor "i" with the kernel for scaling with factor "i+1"
- You must only provide the CUDA code
- Deadline: December 8th