

Final Project OpenMP 1 - KMEANS

Student: Tiago de Souza Oliveira

Toolkit:

- `icc -g -pg -qopenmp -o kmeans_omp kmeans_omp.c`
- `gprof -l kmeans_omp > kmeans_omp_gprof.out`
- `more kmeans_omp_gprof.out`

```
openmp > $ job_kmeans_omp.sh
2  #SBATCH -o %x-%J.out
3  #SBATCH -e %x-%J.error
4
5  #SBATCH -J omp_job_kmeans      # Job name
6  #SBATCH -o omp_job_kmeans.o%j  # Name of stdout output file(%j expands to jobId)
7  #SBATCH -e omp_job_kmeans.o%j  # Name of stderr output file(%j expands to jobId)
8
9  #SBATCH --time=0-00:05:00 #requested time to run the job
10 #SBATCH -c 32 #(64 cores per job)
11 #SBATCH -t 00:10:00 #(10 min of execution time)
12 #SBATCH --mem=16GB #(4GB of memory)
13 #SBATCH --exclusive
14
15 export OMP_NUM_THREADS=8
16 time ./kmeans_omp
```

- For sbatch:

By sticking to the Incremental Parallelization strategy, I tried to focus on the most expensive parts of the application. From that perspective, the gprof report gives some hints as it shows below in terms of routine and subroutine time consumption, computational cost, call frequency and so on.

```
1  Flat profile:
2
3  Each sample counts as 0.01 seconds.
4  % cumulative self self total
5  time seconds seconds calls ps/call ps/call name
6  92.27 415.97 415.97 distEucl (kmeans_omp.c:355 @ 401f8c)
7  5.06 438.79 22.81 distEucl (kmeans_omp.c:354 @ 401fee)
8  1.28 444.54 5.76 recalculateCenters (kmeans_omp.c:322 @ 401d8c)
9  0.72 447.81 3.27 recalculateCenters (kmeans_omp.c:321 @ 401d14)
10 0.18 448.60 0.79 __libm_sqrt_ex
11 0.13 449.18 0.58 findClosestCenters (kmeans_omp.c:292 @ 401b88)
12 0.12 449.73 0.55 createRandomVectors (kmeans_omp.c:151 @ 40143b)
13 0.08 450.09 0.36 recalculateCenters (kmeans_omp.c:320 @ 401dee)
75 granularity: each sample hit covers 2 byte(s) for 0.00% of 451.11 seconds
76
77 index % time self children called name
78
79 [5] 0.2 0.79 0.00 <spontaneous>
80 __libm_sqrt_ex [5]
81
82 [14] 0.0 0.05 0.00 15000000/15000000 findClosestCenters (kmeans_omp.c:292 @ 401b88) [6]
83 distEucl (kmeans_omp.c:349 @ 401f53) [14]
84
85 [25] 0.0 0.00 0.00 150000/150000 findClosestCenters (kmeans_omp.c:293 @ 401bf0) [21]
86 argMin (kmeans_omp.c:368 @ 402019) [25]
87
88 [26] 0.0 0.00 0.00 1/15 kMeans (kmeans_omp.c:188 @ 401669) [69]
89 0.00 0.00 14/15 kMeans (kmeans_omp.c:188 @ 4016a4) [71]
90 0.00 0.00 15 findClosestCenters (kmeans_omp.c:283 @ 401b2e) [26]
```

The gprof report tells essentially that one routine alone consumes more than 95% of the time of processing. The second one relies at a far lower level.

- `distEuclid`
- and `recalculateCenters`


Although I have tried to also improve parallelism for other functions, for most of them it results that the fork-join or thread-startup side effect cost was higher than the speeding up parallelism benefit as the more threads I added to the tests.

When focusing on parallelizing loops inside distEuclid function, I noticed that the more threads I added to the application the lower it becomes. So I realized that the root cause for that could be external loops or upper level loops creating such unnecessary fork-join overhead.

Therefore, by placing the OMP parallelism at the function level before calling distEuclid (findClosestCenters) results in a better balance for speedup and consistency in the application.

Speedup for running in a node with 64 cores per job configuration

Threads OMP (findClosestCenters)					
Threads	2	4	8	16	32
Speedup	1.84	3.53	6.61	11.44	18.30
parallel (s)	299	156	83	48	30
	300	156	83	48	30
	298	155	83	48	30
Average (s)	299	155.67	83	48	30
sequence (s)					
Average (s)	549				



OMP in findClosestCenters function

```

283 double findClosestCenters( double patterns[][Nv], double centers[][Nv], int classes[], double ***distances ) {
284     double error = 0.0 ;
285     size_t i, j ;
286
287     //placing the parallel for here in order to reduce overhead of creating/destroying threads per (i * j) in the nested loop inside functions
289     #pragma omp parallel for private(j) reduction(+=error) schedule(static)
290     for ( i = 0; i < N; i++ ) {
291         for ( j = 0; j < Nc; j++ )
292             (* distances)[i][j] = distEuclid( patterns[i], centers[j] ) ;
293         classes[i] = argMin( (* distances)[i], Nc ) ;
294         error += (* distances)[i][classes[i]] ;
295     }
296
297     return error;
298 }
299


```

Adding OMP also in recalculateCenters function

After adding parallelism to the second most time expensive function(recaculateCenters), the following statistics could be observed. In summary, the speedup could get better as expected with the caveat of correctly placing the barrier to keep the consistency as before for the sequential implementation.

The chart tells us that the application has such linear performance behavior with respect to threads added to speedup.

Speedup for running in a node with 64 cores per job configuration(findClosestCenters, recalculateCenters)



OMP in recalculateCenters function

```

309 void recalculateCenters( double patterns[][Nv], double centers[][Nv], int classes[], double ***y, double ***z ) {
310     double error = 0.0 ;
311     size_t i, j;
312     double t1, t2;
313     t1=omp_get_wtime();
314     #pragma omp parallel for private(j) reduction(+:error)
315     for ( i = 0; i < Nv; i++ ) {
316         for ( j = 0; j < Nv; j++ ) {
317             (* y)[classes[i]][j] += patterns[i][j] ;
318             (* z)[classes[i]][j] ++ ;
319         }
320     }
321     //barrier to make sure that all threads have finished the calculation of tmp arrays
322     #pragma omp barrier
323     // update step of centers
324     for ( i = 0; i < Nv; i++ ) {
325         for ( j = 0; j < Nv; j++ ) {
326             centers[i][j] = (* y)[i][j]/(* z)[i][j] ;
327             (* y)[i][j] = 0.0 ;
328             (* z)[i][j] = 0.0 ;
329         }
330     }
331     t2=omp_get_wtime();
332     printf("\nTime taken for recalculateCenters is %g seconds\n\n", t2-t1);
333     return ;
334 }

```