

Final Project OpenMP 2 - SEISMIC

Student: Tiago de Souza Oliveira

Toolkit:

- `icc -g -pg -qopenmp -o seismic_omp seismic_omp.c`
- `gprof -l seismic_omp > seismic_omp_gprof.out`
- `more seismic_omp_gprof.out`

```
openmp > $ job_seismic_omp.sh

1  #!/bin/bash
2  #SBATCH -o %x-%J.out
3  #SBATCH -e %x-%J.error
4
5  #SBATCH -J omp_job_seismic      # Job name
6  #SBATCH -o omp_job_seismic.o%j  # Name of stdout output file(%j expands to jobId)
7  #SBATCH -e omp_job_seismic.o%j  # Name of stderr output file(%j expands to jobId)
8
9  #SBATCH --time=0-00:05:00 #requested time to run the job
10 #SBATCH -c 32 #(64 cores per job)
11 #SBATCH -t 00:10:00 #(10 min of execution time)
12 #SBATCH --mem=16GB #(4GB of memory)
13 #SBATCH --exclusive
14
15 export OMP_NUM_THREADS=8
16 time ./seismic_omp
```

- For sbatch:

By sticking to the Incremental Parallelization strategy, I tried to focus on the most expensive parts of the application. From that perspective, the gprof report gives some hints as it shows below in terms of routine and subroutine time consumption, computational cost, call frequency and so on.

```
1  Flat profile:
2
3  Each sample counts as 0.01 seconds.
4
5  %   cumulative   self           self       total
6  time   seconds  seconds   calls   Ts/call   Ts/call   name
7  65.77    296.39    296.39                main (seismic_omp.c:321 @ 401fdb)
8  21.53    393.41    97.03                main (seismic_omp.c:320 @ 401fa1)
9  10.01    438.53    45.12                main (seismic_omp.c:319 @ 401f5d)
10  1.04    443.24    4.71                 main (seismic_omp.c:300 @ 401da8)
11  0.57    445.79    2.55                main (seismic_omp.c:298 @ 401d4d)
12  0.18    446.62    0.83                main (seismic_omp.c:301 @ 401dd2)
13  0.14    447.24    0.62                main (seismic_omp.c:275 @ 401b70)
14  0.13    447.83    0.59                main (seismic_omp.c:332 @ 402111)
15  0.12    448.38    0.55                main (seismic_omp.c:296 @ 401cf4)
```

```
72  granularity: each sample hit covers 2 byte(s) for 0.00% of 450.95 seconds
73
74  index % time   self children   called    name
75  -----
76  0.00  0.00    1/151500039  main (seismic_omp.c:165 @ 40134c) [222]
77  0.00  0.00    1/151500039  main (seismic_omp.c:167 @ 401351) [223]
78  0.00  0.00    1/151500039  main (seismic_omp.c:205 @ 4015cb) [241]
79  0.00  0.00    1/151500039  main (seismic_omp.c:206 @ 4015ed) [242]
80  0.00  0.00    1/151500039  main (seismic_omp.c:207 @ 401617) [243]
81  0.00  0.00    1/151500039  main (seismic_omp.c:211 @ 401641) [244]
82  0.00  0.00    1/151500039  main (seismic_omp.c:221 @ 40168f) [250]
83  0.00  0.00    1/151500039  main (seismic_omp.c:228 @ 4016ff) [254]
84  0.00  0.00    1/151500039  main (seismic_omp.c:237 @ 401876) [262]
85  0.00  0.00    2/151500039  main (seismic_omp.c:236 @ 40184c) [261]
86  0.00  0.00   14/151500039  main (seismic_omp.c:226 @ 4016bb) [252]
87  0.00  0.00   14/151500039  main (seismic_omp.c:229 @ 401729) [255]
88  0.00  0.00  1500000/151500039  main (seismic_omp.c:280 @ 401c05) [284]
89  0.00  0.00  150000000/151500039  main (seismic_omp.c:275 @ 401b70) [7]
90  [29]  0.0    0.00  0.00  151500039  main (seismic_omp.c:134 @ 4012e8) [29]
```

It is hard to interpret the regular gprof report. Therefore, I could not come to a conclusion based on those statistics. Nevertheless, I could reason about functions called from the long Main implementation, where I spot **smvp** as a classic vector product algorithm

For that, I performed another query to gprof, but this time only for getting information about the smvp function.

prof -psmvp -b seismic_omp gmon.out > analysis.txt

```

openmp > ≡ analysis.txt
1  Flat profile:
2
3  Each sample counts as 0.01 seconds.
4  %      cumulative      self           self       total
5  time    seconds  seconds    calls  ms/call  ms/call  name
6  100.01      16.50      16.50      3855     4.28     4.28    smvp
7

```

In order to better evaluate the parallelism approach that better fits from the efficiency and consistency perspectives, I plan to setup an OMP environment on my ubuntu machine with appropriate profiling tools like Intel Advisor, Valgrind/Callgrind, CloverLeaf/KCacheGrind and some other gprof parameters as well. Then I think I would be able to better design the parallelism, which I have been thinking would demand some changes in the code..

To do so, I have played the hiphotesis that if I could set the accumulator apart from the array and use a local variable instead, I would be able to leverage “**#pragma omp parallel firstprivate (Anext) reduction (+ : variable1,...)**”, for instance. That's the incomplete implementation that can be seen in the code.

The tests show inconsistencies when calculating **epicenternode** values.

```

1212 double sum0, sum1, sum2;
1213 double sumw0, sumw1, sumw2;
1214
1215 for (i = 0; i < nodes; i++) {
1216     Anext = Aindex[i];
1217     Alast = Aindex[i + 1];
1218
1219     sum0 = A[Anext][0][0]*v[i][0] + A[Anext][0][1]*v[i][1] + A[Anext][0][2]*v[i][2];
1220     sum1 = A[Anext][1][0]*v[i][0] + A[Anext][1][1]*v[i][1] + A[Anext][1][2]*v[i][2];
1221     sum2 = A[Anext][2][0]*v[i][0] + A[Anext][2][1]*v[i][1] + A[Anext][2][2]*v[i][2];
1222
1223     sumw0 = w[Acol[Anext + 1]][0], sumw1 = w[Acol[Anext + 1]][1], sumw2 = w[Acol[Anext + 1]][2];
1224
1225 //Hypothesis that replacing array with a local variable, to use the cache better, avoid false sharing and improve locality
1226 #pragma omp parallel firstprivate ( Anext ) reduction ( + : sum0, sum1, sum2, sumw0, sumw1, sumw2 )
1227 {
1228     Anext++;
1229     while (Anext < Alast) {
1230         col = Acol[Anext];
1231
1232         sum0 += A[Anext][0][0]*v[col][0] + A[Anext][0][1]*v[col][1] + A[Anext][0][2]*v[col][2];
1233         sum1 += A[Anext][1][0]*v[col][0] + A[Anext][1][1]*v[col][1] + A[Anext][1][2]*v[col][2];
1234         sum2 += A[Anext][2][0]*v[col][0] + A[Anext][2][1]*v[col][1] + A[Anext][2][2]*v[col][2];
1235
1236         /*
1237         w[col][0] += A[Anext][0][0]*v[i][0] + A[Anext][1][0]*v[i][1] + A[Anext][2][0]*v[i][2];
1238         w[col][1] += A[Anext][0][1]*v[i][0] + A[Anext][1][1]*v[i][1] + A[Anext][2][1]*v[i][2];
1239         w[col][2] += A[Anext][0][2]*v[i][0] + A[Anext][1][2]*v[i][1] + A[Anext][2][2]*v[i][2];
1240         */
1241
1242         sumw0 += w[col][0] + A[Anext][0][0]*v[i][0] + A[Anext][1][0]*v[i][1] + A[Anext][2][0]*v[i][2];
1243         sumw1 += w[col][1] + A[Anext][0][1]*v[i][0] + A[Anext][1][1]*v[i][1] + A[Anext][2][1]*v[i][2];
1244         sumw2 += w[col][2] + A[Anext][0][2]*v[i][0] + A[Anext][1][2]*v[i][1] + A[Anext][2][2]*v[i][2];
1245
1246         Anext++;
1247     }
1248
1249     w[i][0] = sumw0 + sum0;
1250     w[i][1] = sumw1 + sum1;
1251     w[i][2] = sumw2 + sum2;
1252 }
1253
1254 }
1255

```

In summary, this is still a Work In Progress for me, and I should invest more time on a better profiling tool for a loop interaction to choose the right OMP parallelism instructions and the correct changes in the code to accommodate higher scalability.